

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import StandardScaler
import seaborn as sns
```

## Reading the diabetes dataset

(Downloaded from [here \(https://www.kaggle.com/uciml/pima-indians-diabetes-database\)](https://www.kaggle.com/uciml/pima-indians-diabetes-database).)

```
In [ ]: #reading the diabetes dataset from the folder and storing i
t in a dataframe named diabetes
diabetes = pd.read_csv("./diabetes.csv")
```

## Dimensions of the dataframe

```
In [ ]: diabetes.shape
```

Out[ ]: (768, 9)

## Sample data in dataframe

```
In [ ]: diabetes.head()
```

Out[ ]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPe
0	6	148	72	35	0	33.6	
1	1	85	66	29	0	26.6	
2	8	183	64	0	0	23.3	
3	1	89	66	23	94	28.1	
4	0	137	40	35	168	43.1	

## Problem description:

- *Binary classification problem*
- *Label : 'Outcome' = 0 or 1*
- *Features :*
  - 'Pregnancies'
  - 'Glucose'
  - 'BloodPressure'
  - 'SkinThickness'
  - 'Insulin'
  - 'BMI'
  - 'DiabetesPedigreeFunction'
  - 'Age'

## Summary of the data


```
In [ ]: diabetes.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Pregnancies                          768 non-null    int64
1   Glucose                              768 non-null    int64
2   BloodPressure                        768 non-null    int64
3   SkinThickness                       768 non-null    int64
4   Insulin                             768 non-null    int64
5   BMI                                  768 non-null    float64
6   DiabetesPedigreeFunction             768 non-null    float64
7   Age                                  768 non-null    int64
8   Outcome                              768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

## Step 1: Visualisation (part 1)

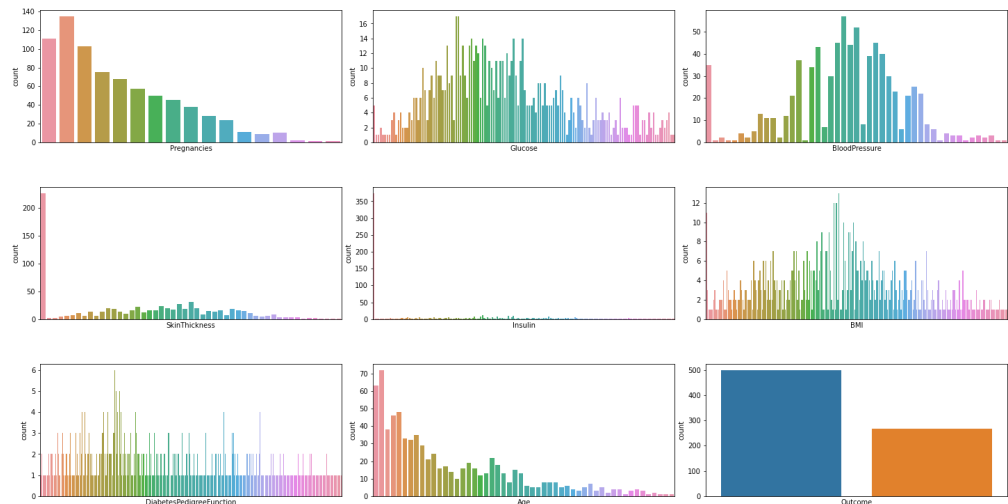
Since we have discrete categorical data, we plot the estimate plots using countplot

```
In [ ]: sns.countplot(diabetes['Pregnancies']);
```



### Plotting the rest of the features using countplots

```
In [ ]: fig, axs = plt.subplots(ncols=3, nrows=3, figsize=(20, 10))
index = 0
axs = axs.flatten()
for i in diabetes.columns:
    sns.countplot(diabetes[i], ax=axs[index])
    axs[index].set_xticks([])
    index+=1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```



**Inference:**

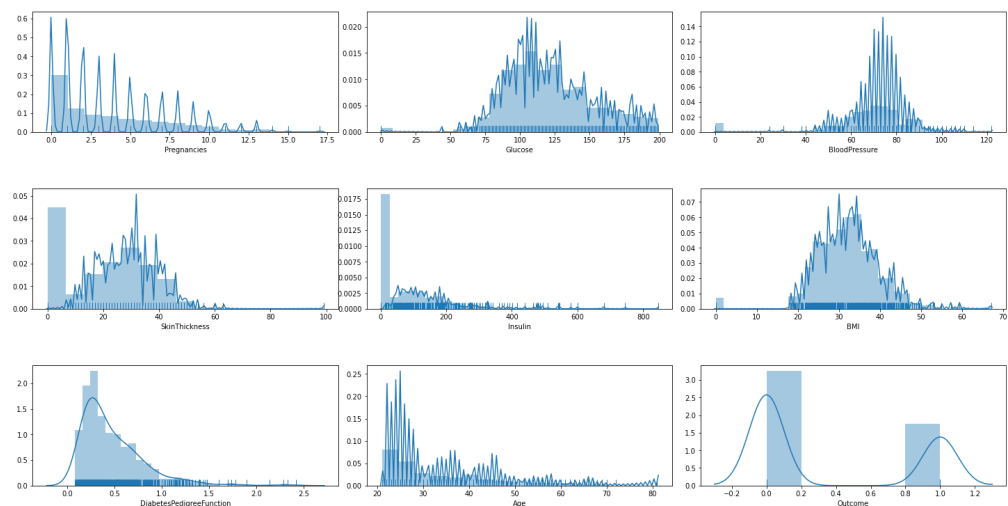
- More data of non-diabetics present than diabetics.
- Prediction accuracy whether a person is non-diabetic will be higher than him/her being diabetic.
- For features Skin Thickness, and Insulin, zero values must be imputed.

**Visualising using distplot : evenness of data spread**

**Output:** A curve that roughly fits the distribution.

(We also add a rugplot which marks each individual point on the x-axis)

```
In [ ]: fig, axs = plt.subplots(ncols=3, nrows=3, figsize=(20, 10))
index = 0
axs = axs.flatten()
for k,v in diabetes.items():
    sns.distplot(v, ax=axs[index], kde_kws={'bw': 0.1}, rug=True)
    index += 1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```

**Inference:**

- Glucose, BP, BMI have approximately normally distributed data.
- DiabetesPedigreeFunction peak has a slight shift to the left.
- SkinThickness, Insulin have a sharp spike, due to imputation of 0 values with a single fixed values.

To handle this, we use normalization/ standardisation further.

## Step 2: Checking for missing value and string datatype and abnormal values

```
In [ ]: diabetes.isnull().sum()
```

```
Out[ ]: Pregnancies      0
        Glucose          0
        BloodPressure    0
        SkinThickness    0
        Insulin          0
        BMI              0
        DiabetesPedigreeFunction  0
        Age              0
        Outcome          0
        dtype: int64
```

**Inference:** *No null data present*

```
In [ ]: diabetes.dtypes
```

```
Out[ ]: Pregnancies      int64
        Glucose          int64
        BloodPressure    int64
        SkinThickness    int64
        Insulin          int64
        BMI              float64
        DiabetesPedigreeFunction  float64
        Age              int64
        Outcome          int64
        dtype: object
```

**Inference:** *No string or object datatype present*

```
In [ ]: pd.DataFrame(diabetes[:,]==0).sum()
```

```
Out[ ]: Pregnancies      111
        Glucose          5
        BloodPressure    35
        SkinThickness    227
        Insulin          374
        BMI              11
        DiabetesPedigreeFunction  0
        Age              0
        Outcome          500
        dtype: int64
```

**Inference:** *Glucose, BP, SkinThickness, Insulin, BMI can't be 0 - data has to be processed*

### Step 3: Imputing the abnormal values

```
In [ ]: diabetes1 = diabetes.copy(deep=True) # copy of dataframe made in order to keep original dataframe unchanged
```

Using mean to impute the zero values for columns: Glucose, BP, SkinThickness, Insulin, BMI

```
In [ ]: columns = ['Glucose', 'BloodPressure', 'SkinThickness', 'BMI']

for i in columns:
    avg = diabetes1[i][diabetes1[i]>0].mean()
    diabetes1[i] = diabetes1[i].replace(to_replace=0, value=avg)
```

```
In [ ]: md = diabetes1['Insulin'][diabetes1['Insulin']>0].mode()[0]
diabetes1['Insulin'] = diabetes1['Insulin'].replace(to_replace=0, value=md)
```

```
In [ ]: pd.DataFrame(diabetes1[:,]==0).sum()
```

```
Out[ ]: Pregnancies      111
Glucose                0
BloodPressure          0
SkinThickness          0
Insulin                0
BMI                    0
DiabetesPedigreeFunction  0
Age                    0
Outcome                500
dtype: int64
```

**Inference:** All zero values replaced with mean of the rest of the values.

**Possible difficulty:** The distribution of data maybe spiked since there were lot of zero values in particularly 'SkinThickness' and 'Insulin' columns.

### Step 4: Data analysis and visualisation(part 2)

```
In [ ]: diabetes1.describe()
```

```
Out[ ]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI
<b>count</b>	768.000000	768.000000	768.000000	768.000000	768.000000	768.0
<b>mean</b>	3.845052	121.686763	72.405184	29.153420	130.932292	32.4
<b>std</b>	3.369578	30.435949	12.096346	8.790942	88.700443	6.8
<b>min</b>	0.000000	44.000000	24.000000	7.000000	14.000000	18.2
<b>25%</b>	1.000000	99.750000	64.000000	25.000000	105.000000	27.5
<b>50%</b>	3.000000	117.000000	72.202592	29.153420	105.000000	32.4
<b>75%</b>	6.000000	140.250000	80.000000	32.000000	127.250000	36.6
<b>max</b>	17.000000	199.000000	122.000000	99.000000	846.000000	67.1

**Inference:** High variation in all columns


**Possible solution:** Scaling of data: either normalization or standardization

```
In [ ]: fig, axs = plt.subplots(ncols=9, nrows=1, figsize=(20, 5))
index = 0
axs = axs.flatten()

# for k,v in diabetes.items():
#     sns.boxplot(y=v, data=diabetes, ax=axs[index])
#     index += 1

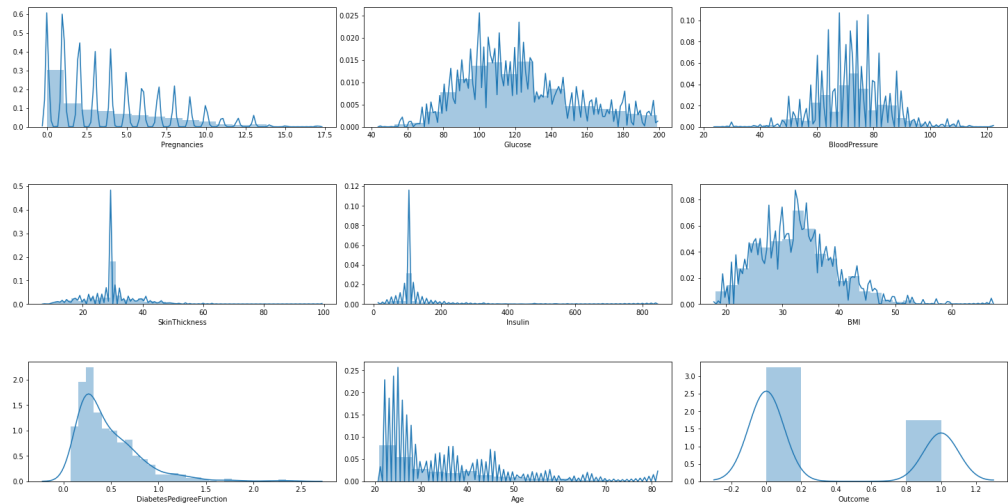
for k,v in diabetes1.items():
    sns.boxplot(y=v, data=diabetes1, ax=axs[index])
    index += 1

plt.tight_layout(pad=0.4, w_pad=0.1, h_pad=5.0)
```



**Inference:** Imputation of abnormal values and plotting the boxplot, shows outliers towards the lower range have been successfully removed.

```
In [ ]: fig, axs = plt.subplots(ncols=3, nrows=3, figsize=(20, 10))
index = 0
axs = axs.flatten()
for k,v in diabetes1.items():
    sns.distplot(v, ax=axs[index], kde_kws={'bw': 0.1}) # for some prob write kde
    index += 1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```



**Inference:** Imputation of abnormal values and plotting the distplot, shows outliers towards the lower range have been successfully removed.

## Step 5: Treating outliers

```
In [ ]: diabetes1_0 = diabetes1.copy(deep=True)
```

```
In [ ]: for i in diabetes1_0.columns:
    upper = diabetes1_0[i].mean() + diabetes1_0[i].std()*
    3.1
    print(upper)
    diabetes1_0 = diabetes1_0[~(diabetes1_0[i] >= upper)]
```

```
14.290744077699808
215.9613909365196
109.95551519809001
56.38625954346237
407.2634508820364
52.81267394865042
1.4346189055800058
69.44588978318626
1.8024381326507237
```




```
In [ ]: diabetes1_0.shape
```

```
Out[ ]: (722, 9)
```

```
In [ ]: fig, axs = plt.subplots(ncols=9, nrows=1, figsize=(20, 5))
        index = 0
        axs = axs.flatten()

        for k,v in diabetes1_0.items():
            sns.boxplot(y=v, data=diabetes1_0, ax=axs[index])
            index += 1

        plt.tight_layout(pad=0.4, w_pad=0.1, h_pad=5.0)
```



### Inference :

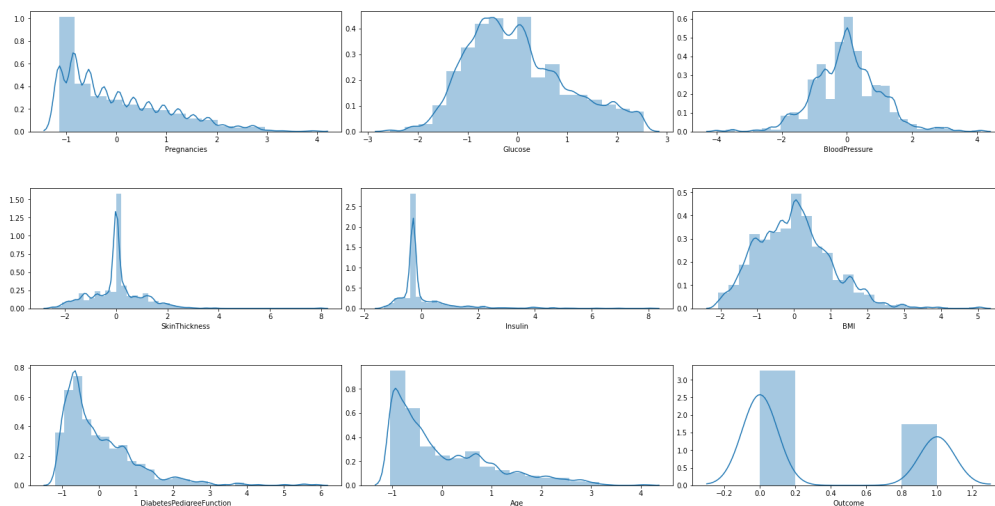
Manual outlier not preferable as automatic outlier detection (see train test part) gives better accuracy score for ML models.

## Step 6: Normalization/ Standardisation of data (visualisation part-2)

```
In [ ]: diabetes2 = diabetes1.copy(deep=True)
        sc = StandardScaler()
```

```
In [ ]: # scaling all columns except 1st and last
        for i in diabetes2.columns[0:-1]:
            diabetes2[i] = sc.fit_transform(pd.DataFrame(diabetes2.
                loc[:, i]).values)
```

```
In [ ]: fig, axs = plt.subplots(ncols=3, nrows=3, figsize=(20, 10))
index = 0
axs = axs.flatten()
for k,v in diabetes2.items():
    sns.distplot(v, ax=axs[index], kde_kws={'bw': 0.1}) # for some prob write kde
    index += 1
plt.tight_layout(pad=0.4, w_pad=0.5, h_pad=5.0)
```



### Inference:

- Variation is less now; with mean as 0 and standard deviation as 1
- Though the peaks in SkinThickness and Insulin could not be removed, the desired normal distribution has been achieved.

```
In [ ]: diabetes2.describe()
```

```
Out[ ]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin
<b>count</b>	7.680000e+02	7.680000e+02	7.680000e+02	7.680000e+02	7.680000e+02
<b>mean</b>	2.544261e-17	-3.301757e-16	6.966722e-16	6.866252e-16	3.122502e-17
<b>std</b>	1.000652e+00	1.000652e+00	1.000652e+00	1.000652e+00	1.000652e+00
<b>min</b>	-1.141852e+00	-2.554131e+00	-4.004245e+00	-2.521670e+00	-1.319142e+00
<b>25%</b>	-8.448851e-01	-7.212214e-01	-6.953060e-01	-4.727737e-01	-2.925486e-01
<b>50%</b>	-2.509521e-01	-1.540881e-01	-1.675912e-02	8.087936e-16	-2.925486e-01
<b>75%</b>	6.399473e-01	6.103090e-01	6.282695e-01	3.240194e-01	-4.154084e-02
<b>max</b>	3.906578e+00	2.541850e+00	4.102655e+00	7.950467e+00	8.066856e+00

**Inference:** Variation is less now.

## Step 7: Exploring linearity of data (visualisation - part 3)

In [ ]: `diabetes2.corr()`

Out[ ]:

	Pregnancies	Glucose	BloodPressure	SkinThickness
Pregnancies	1.000000	0.127911	0.208522	0.082989
Glucose	0.127911	1.000000	0.218367	0.192991
BloodPressure	0.208522	0.218367	1.000000	0.192816
SkinThickness	0.082989	0.192991	0.192816	1.000000
Insulin	0.005204	0.411642	0.027149	0.150020
BMI	0.021565	0.230941	0.281268	0.542398
DiabetesPedigreeFunction	-0.033523	0.137060	-0.002763	0.100966
Age	0.544341	0.266534	0.324595	0.127872
Outcome	0.221898	0.492928	0.166074	0.215299

**Inference:** *Comparatively higher correlation between*


- *age and pregnancies, which is normal.*
- *skin thick and BMI, which also can be related*

*Here maximum correlation is 0.54.*

*Had there been any correlation value been > 0.8 we would have selected one feature of the two correlated feature.*

*Here we are unable to eliminate any features.*

```
In [ ]: fig, ax = plt.subplots(figsize=(10,7))
sns.heatmap(diabetes2.corr(), annot = True, cmap="coolwarm");
```



**Inference:** *Comparatively higher correlation between*

- Age and pregnancies, which is normal.
- SkinThickness and BMI, which also can be related

had there been any correlation value been  $> 0.8$  we would have selected one feature of the two

## Step 8: Machine Learning Models

This is a **binary classification** problem.

The models we will apply are:

1. Logistic Regression
2. Naive Bayes Classifier (Gaussian)
3. SVM (linear, poly, radial kernel)
4. Decision Tree
5. Random Forest

Then we conclude which is the best model to be applied.

## Importing the Libraries

```
In [ ]: from sklearn.linear_model import LogisticRegression
        from sklearn.naive_bayes import GaussianNB
        from sklearn import svm
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.model_selection import train_test_split, KFold, cross_val_score
        from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
```

## Pre-step 1: Diving into train and test set

```
In [ ]: X = diabetes2.iloc[:, :8].values
```

```
In [ ]: Y = diabetes2.iloc[:,8].values
```

```
In [ ]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.1, random_state=1)
```

```
In [ ]: # kfold = KFold(n_splits=10, random_state=10)
```

## Pre-step 2: automatic outlier detection

```
In [ ]: from sklearn.svm import OneClassSVM
        ee = OneClassSVM(nu=0.01)
        yhat = ee.fit_predict(X_train)
        # select all rows that are not outliers
        mask = yhat != -1
        X_train, Y_train = X_train[mask, :], Y_train[mask]
```

## Setting DataFrame to store accuracies

```
In [ ]: acc_stats = pd.DataFrame(columns=['Algorithm used', 'Train Score', 'Test Score'])
```

## Model 1: Logistic Regression

```
In [ ]: log = LogisticRegression()
```

```
In [ ]: log.fit(X_train, Y_train)
```

```
Out[ ]: LogisticRegression(C=1.0, class_weight=None, dual=False, fit_intercept=True,
                             intercept_scaling=1, l1_ratio=None, max_iter=100,
                             multi_class='auto', n_jobs=None, penalty='l2',
                             random_state=None, solver='lbfgs', tol=0.0001, verbose=0,
                             warm_start=False)
```

```
In [ ]: Y_pred = log.predict(X_test)
```

```
In [ ]: Y_pred
```

```
Out[ ]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0,
                0, 0, 0, 0, 0,
                1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0,
                1, 0, 0, 0, 0,
                0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1,
                0, 1, 0, 0, 0,
                1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1])
```

```
In [ ]: Y_test
```

```
Out[ ]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1,
                0, 0, 0, 1, 1,
                1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 0, 0, 0,
                0, 0, 0, 0, 1,
                0, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0,
                0, 1, 0, 1, 0,
                1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1])
```

```
In [ ]: tstSc = accuracy_score(Y_pred, Y_test) # store test score
tstSc
```

```
Out[ ]: 0.7662337662337663
```

```
In [ ]: trSc = log.score(X_train, Y_train) # store train score
trSc
```

```
Out[ ]: 0.7791411042944786
```

```
In [ ]: confusion_matrix(Y_pred, Y_test)
```

```
Out[ ]: array([[41, 11],
               [ 7, 18]])
```

```
In [ ]: print(classification_report(Y_pred, Y_test))
```

	precision	recall	f1-score	support
0	0.85	0.79	0.82	52
1	0.62	0.72	0.67	25
accuracy			0.77	77
macro avg	0.74	0.75	0.74	77
weighted avg	0.78	0.77	0.77	77

```
In [ ]: # cross_val_score(log, X, Y, cv=kfold, scoring='accuracy').
        mean()
```

```
In [ ]: temp = pd.DataFrame([["Logistic Regression", trSc, tstSc]],
                             columns=['Algorithm used', 'Train Score', 'Test Score'])
acc_stats = pd.concat([acc_stats, temp], sort=False, ignore_index=True)
```

#### Inference:

Train Score : 75

Test Score : 77

#### Model 2: Naive Bayes (Gaussian)

```
In [ ]: nvclassifier = GaussianNB()
nvclassifier.fit(X_train, Y_train)
y_pred = nvclassifier.predict(X_test)
tstSc = accuracy_score(y_pred, Y_test)
trSc = nvclassifier.score(X_train, Y_train)
print("Test Score: ", tstSc, "\nTrain Score: ", trSc)
```

Test Score: 0.7662337662337663  
Train Score: 0.7484662576687117

```
In [ ]: confusion_matrix(y_pred, Y_test)
```

```
Out[ ]: array([[41, 11],
               [ 7, 18]])
```

```
In [ ]: print(classification_report(y_pred, Y_test))
```

	precision	recall	f1-score	support
0	0.85	0.79	0.82	52
1	0.62	0.72	0.67	25
accuracy			0.77	77
macro avg	0.74	0.75	0.74	77
weighted avg	0.78	0.77	0.77	77

```
In [ ]: temp = pd.DataFrame([["Naive Bayes Classifier (Gaussian)",
trSc, tstSc]], columns=['Algorithm used', 'Train Score', 'T
est Score'])
acc_stats = pd.concat([acc_stats, temp], sort=False, ignore
_index=True)
```

### Model 3: SVM

#### Using Linear Kernel

```
In [ ]: clf = svm.SVC(kernel='linear')
clf.fit(X_train, Y_train)
y_pred = clf.predict(X_test)
tstSc = accuracy_score(y_pred, Y_test)
trSc = clf.score(X_train, Y_train)
print("Test Score: ", tstSc, "\nTrain Score: ", trSc)
```

```
Test Score:  0.7922077922077922
Train Score:  0.7714723926380368
```

```
In [ ]: temp = pd.DataFrame([["SVM (Linear Kernel)", trSc, tstSc]],
columns=['Algorithm used', 'Train Score', 'Test Score'])
acc_stats = pd.concat([acc_stats, temp], sort=False, ignore
_index=True)
```

**Inference:** 76 76

#### Using Polynomial Kernel



```
In [ ]: clf = svm.SVC(kernel='poly', C=0.1)
        clf.fit(X_train, Y_train)
        y_pred = clf.predict(X_test)
        tstSc = accuracy_score(y_pred, Y_test)
        trSc = clf.score(X_train, Y_train)
        print("Test Score: ", tstSc, "\nTrain Score: ", trSc)
```

```
Test Score:  0.6883116883116883
Train Score:  0.74079754601227
```

```
In [ ]: temp = pd.DataFrame([["SVM (Polynomial Kernel)", trSc, tstSc],
                             columns=['Algorithm used', 'Train Score', 'Test Score']])
        acc_stats = pd.concat([acc_stats, temp], sort=False, ignore_index=True)
```

### Using Radial Kernel

```
In [ ]: clf = svm.SVC(kernel='rbf')
        clf.fit(X_train, Y_train)
        y_pred = clf.predict(X_test)
        tstSc = accuracy_score(y_pred, Y_test)
        trSc = clf.score(X_train, Y_train)
        print("Test Score: ", tstSc, "\nTrain Score: ", trSc)
```

```
Test Score:  0.8051948051948052
Train Score:  0.8205521472392638
```

```
In [ ]: temp = pd.DataFrame([["SVM (Radial Kernel)", trSc, tstSc]],
                             columns=['Algorithm used', 'Train Score', 'Test Score'])
        acc_stats = pd.concat([acc_stats, temp], sort=False, ignore_index=True)
```

```
In [ ]: print("Confusion matrix:\n", confusion_matrix(y_pred, Y_test))
        print(classification_report(y_pred, Y_test))
```

Confusion matrix:

```
[[45 12]
 [ 3 17]]
```

	precision	recall	f1-score	support
0	0.94	0.79	0.86	57
1	0.59	0.85	0.69	20
accuracy			0.81	77
macro avg	0.76	0.82	0.78	77
weighted avg	0.85	0.81	0.81	77

**Inference:** This model has the best accuracy, as we have compared later. Analysis of Confusion matrix:

- Precision and recall values of 0 (no diabetes) is significantly more than 1.  
Thus, the model can make better prediction that a person does **NOT have Diabetes**.

## Model 4: Decision Tree

```
In [ ]: model=DecisionTreeClassifier(criterion='entropy',splitter='
best',random_state=1, min_samples_split=0.1)
model.fit(X_train,Y_train)
```

```
Out[ ]: DecisionTreeClassifier(ccp_alpha=0.0, class_weight=None, cr
iterion='entropy',
                                max_depth=None, max_features=None, m
ax_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impur
ity_split=None,
                                min_samples_leaf=1, min_samples_spli
t=0.1,
                                min_weight_fraction_leaf=0.0, presor
t='deprecated',
                                random_state=1, splitter='best')
```

```
In [ ]: y_pred=model.predict(X_test)
tstSc = accuracy_score(y_pred, Y_test)
trSc = model.score(X_train, Y_train)
print("Test Score: ", tstSc, "\nTrain Score: ", trSc)
print("Conf matirx:\n", confusion_matrix(y_pred, Y_test))
print(classification_report(y_pred, Y_test))
```

Test Score: 0.8051948051948052

Train Score: 0.8128834355828221

Conf matirx:

```
[[41  8]
```

```
 [ 7 21]]
```

	precision	recall	f1-score	support
0	0.85	0.84	0.85	49
1	0.72	0.75	0.74	28
accuracy			0.81	77
macro avg	0.79	0.79	0.79	77
weighted avg	0.81	0.81	0.81	77

```
In [ ]: temp = pd.DataFrame([["Decision Tree", trSc, tstSc]], columns=['Algorithm used', 'Train Score', 'Test Score'])
acc_stats = pd.concat([acc_stats, temp], sort=False, ignore_index=True)
```

## Model 5: Random Forest

```
In [ ]: model=RandomForestClassifier(n_estimators=150,criterion='entropy',random_state=1, min_samples_split=0.1)
model.fit(X_train,Y_train)
y_pred=model.predict(X_test)
tstSc = accuracy_score(y_pred, Y_test)
trSc = model.score(X_train, Y_train)
print("Test Score: ", tstSc, "\nTrain Score: ", trSc)
print("Conf matirx:\n", confusion_matrix(y_pred, Y_test))
print(classification_report(y_pred, Y_test))
```

Test Score: 0.7922077922077922

Train Score: 0.8236196319018405

Conf matirx:

[[44 12]

[ 4 17]]

	precision	recall	f1-score	support
0	0.92	0.79	0.85	56
1	0.59	0.81	0.68	21
accuracy			0.79	77
macro avg	0.75	0.80	0.76	77
weighted avg	0.83	0.79	0.80	77

```
In [ ]: temp = pd.DataFrame([["Random Forest", trSc, tstSc]], columns=['Algorithm used', 'Train Score', 'Test Score'])
acc_stats = pd.concat([acc_stats, temp], sort=False, ignore_index=True)
```

## Conclusion

*Comparing accuracies to select the best model*

```
In [ ]: acc_stats
```

```
Out[ ]:
```

	Algorithm used	Train Score	Test Score
0	Logistic Regression	0.779141	0.766234
1	Naive Bayes Classifier (Gaussian)	0.748466	0.766234
2	SVM (Linear Kernel)	0.771472	0.792208
3	SVM (Polynomial Kernel)	0.740798	0.688312
4	SVM (Radial Kernel)	0.820552	0.805195
5	Decision Tree	0.812883	0.805195
6	Random Forest	0.823620	0.792208

## Conclusion:

In NaiveBayes and SVM(linear) we are getting underfitting, so we do not consider those two. From remaining, we can see SVM(radial) and DecTree have best Test Score. Among those two, SVM(Radial) has slightly better score. Therefore,

**Best model: SVM (Radial kernel)**

**Achieved accuracy: 82.05%**

However, if we look at the classification reports the best balanced model is **Decision Tree**

**Achieved accuracy: 81.28%**

with **f1 score**:

- 0=> 85
- 1 =>74

## Step 9: Further analysis on the data and visualisation

### Clustering - Unsupervised Learning

```
In [ ]: from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters = 2, random_state = 0)
```

```
In [ ]: kmeans.fit(X)
```

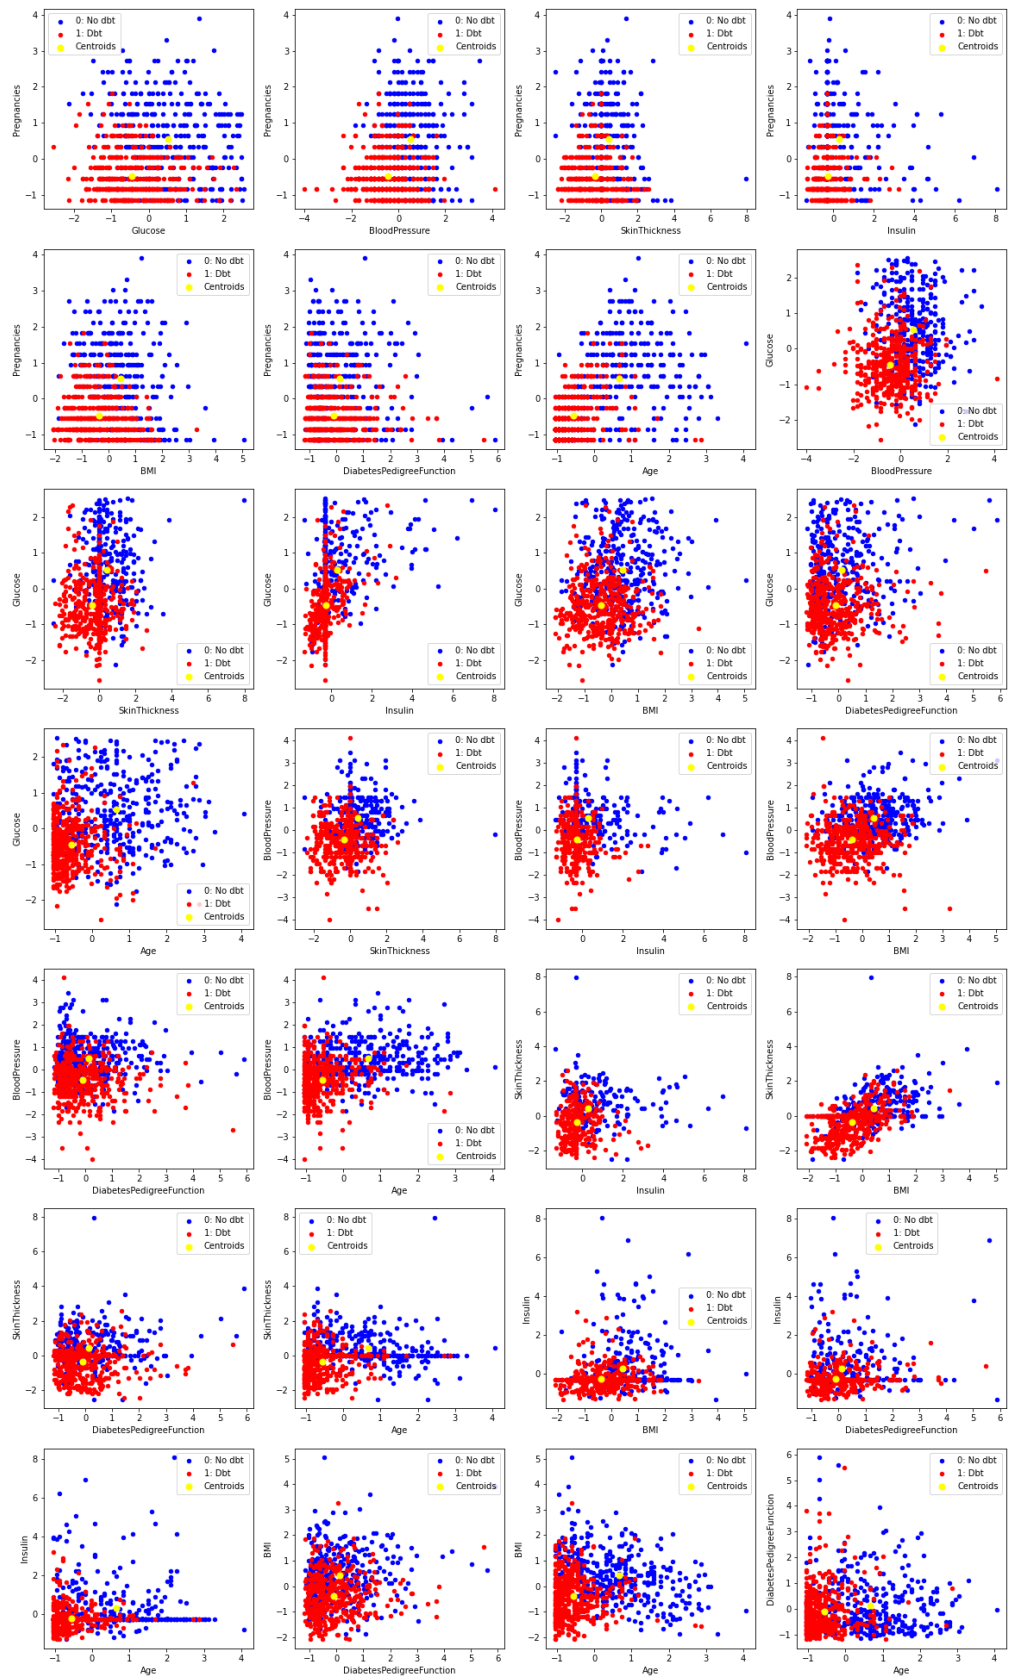
```
Out[ ]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
               n_clusters=2, n_init=10, n_jobs=None, precompute_distances='auto',
               random_state=0, tol=0.0001, verbose=0)
```

```
In [ ]: y_pred = kmeans.predict(X)
```

```
In [ ]: # y_pred
```

```
In [ ]: # Y
```

```
In [ ]: fig, axs = plt.subplots(ncols=4, nrows=7, figsize=(20, 35))
index = 0
axs = axs.flatten()
for i in range(len(X[0])):
    for j in range(i+1, len(X[0])):
        axs[index].scatter(X[y_pred == 0, j], X[y_pred ==
0, i], s = 20, c = 'blue', label = '0: No dbt')
        axs[index].scatter(X[y_pred == 1, j], X[y_pred ==
1, i], s = 20, c = 'red', label = '1: Dbt')
        axs[index].scatter(kmeans.cluster_centers_[:,j], km
eans.cluster_centers_[:,i], s = 50, c = 'yellow', label = '
Centroids')
        axs[index].legend()
        axs[index].set_xlabel(diabetes2.columns[j])
        axs[index].set_ylabel(diabetes2.columns[i])
        index+=1
```



**Inference:**

From the scatter plot of clustering we can conclude that if we are given values for following pair of features we can cluster the data into 2 groups (Diabetes and No diabetes):

- Glucose-Age
- Blood Pressure-Age
- Pregnancies-Age
- Pregnancies-Blood Pressure

**Age and Diabetes**

```
In [ ]: pd.DataFrame(diabetes1.groupby(['Age', 'Outcome'])['Outcome']
                    .count())
```

Out[ ]:

		Outcome
Age	Outcome	
21	0	58
	1	5
22	0	61
	1	11
23	0	31
	1	0
...	...	...
68	0	1
69	0	2
70	1	1
72	0	1
81	0	1

96 rows × 1 columns

**Inference:** The age has several discrete values for that we can divide it into ranges.

Automatic dividing age into range is done using cut

```
In [ ]: diabetes3 = diabetes1.copy(deep=True)
        diabetes3['AgeBand'] = pd.cut(diabetes1['Age'], 8)
```



```
In [ ]: diabetes3[['AgeBand', 'Outcome']].groupby('AgeBand', as_index=False).agg({'Outcome': ['sum', 'count']})
# how many have diabetes in that range
```

Out[ ]:

	AgeBand	Outcome	
		sum	count
0	(20.94, 28.5]	71	367
1	(28.5, 36.0]	70	147
2	(36.0, 43.5]	56	113
3	(43.5, 51.0]	38	68
4	(51.0, 58.5]	22	38
5	(58.5, 66.0]	9	26
6	(66.0, 73.5]	2	8
7	(73.5, 81.0]	0	1

```
In [ ]: diabetes3['AgeBand'] = 0
```

*Age is divided manually into age bands taking 10 years as range.  
For easier realisation and interpretation of data*

```
In [ ]: diabetes3.loc[diabetes3['Age'] <= 30, 'AgeBand'] = 0
diabetes3.loc[(diabetes3['Age'] > 30) & (diabetes3['Age'] <= 40), 'AgeBand'] = 1
diabetes3.loc[(diabetes3['Age'] > 40) & (diabetes3['Age'] <= 50), 'AgeBand'] = 2
diabetes3.loc[(diabetes3['Age'] > 50) & (diabetes3['Age'] <= 60), 'AgeBand'] = 3
diabetes3.loc[diabetes3['Age'] > 60, 'AgeBand'] = 4
```

```
In [ ]: diabetes3
```

```
Out[ ]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Diabetes
0	6	148.0	72.0	35.00000	105	33.6	
1	1	85.0	66.0	29.00000	105	26.6	
2	8	183.0	64.0	29.15342	105	23.3	
3	1	89.0	66.0	23.00000	94	28.1	
4	0	137.0	40.0	35.00000	168	43.1	
...	...	...	...	...	...	...	...
763	10	101.0	76.0	48.00000	180	32.9	
764	2	122.0	70.0	27.00000	105	36.8	
765	5	121.0	72.0	23.00000	112	26.2	
766	1	126.0	60.0	29.15342	105	30.1	
767	1	93.0	70.0	31.00000	105	30.4	

768 rows × 10 columns

```
In [ ]: ageb_outcome_count = diabetes3.groupby(['AgeBand', 'Outcome'])['Outcome'].count()
```

```
In [ ]: ageb_count = diabetes3.groupby(['AgeBand'])['Outcome'].count()
```

```
In [ ]: ageb_analysis = ageb_outcome_count.div(ageb_count, level='AgeBand') * 100
ageb_analysis
```

```
Out[ ]: AgeBand Outcome
0      0      78.417266
      1      21.582734
1      0      51.592357
      1      48.407643
2      0      43.362832
      1      56.637168
3      0      42.592593
      1      57.407407
4      0      74.074074
      1      25.925926
Name: Outcome, dtype: float64
```

```
In [ ]: ageb_analysis[:,1]
```

```
Out[ ]: AgeBand
0      21.582734
1      48.407643
2      56.637168
3      57.407407
4      25.925926
Name: Outcome, dtype: float64
```

**Inference:**

*On comparing the relative percentages of people having diabetes in the different age groupps, we can conclude that:*

*Most prone age group is group 3 that corresponds to (40, 50]*